

CSc 484

# Database Management Systems

Ken Gamradt

Spring 2024

Set Operations – Database Modifications

## Set operations in SQL

- **UNION, INTERSECT, EXCEPT** can be used to combine the results of two or more queries into a single result table
  - **UNION** of two tables A and B, is a table containing all rows that are either in A or B or both
  - **INTERSECT** of two tables A and B, is a table containing all rows that are common to A and B
  - **EXCEPT** of two tables A and B, is a table containing all rows that are in A but not in B



# Set operations in SQL

- Restrictions on using set operations
  - Two tables must contain the same number of columns
    - Their corresponding columns must have the same data types and length
  - It's user's responsibility to ensure that the data values in corresponding columns come from the same domain
    - E.g.: instructor's name and building name

# Set operations in SQL

- The list of all courses taught in Fall 2017

```
select course_id
  from section
 where semester = 'Fall' and year = 2017
```

course\_id

abc Filter...

CS-101

CS-347

PHY-101

- The list of all courses taught in Spring 2018

```
select course_id
  from section
 where semester = 'Spring' and year = 2018
```

course\_id

abc Filter...

CS-101

CS-315

CS-319

CS-319

FIN-201

HIS-351

MU-199



## Set operations in SQL – Union

- **UNION** of two tables A and B, is a table containing all rows that are either in A or B or both
- E.g., find all courses that taught either in Fall 2017 or in Spring 2018

```
(select course_id
  from section
  where semester = 'Fall' and year = 2017)
union  -- key word union between two statements
(select course_id
  from section
  where semester = 'Spring' and year = 2018)
```

course\_id

abc Filter...

CS-101

CS-315

CS-319

CS-347

FIN-201

HIS-351

MU-199

PHY-101

## Set operations in SQL – Union

- The **UNION** operation automatically eliminates duplicates

	course_id
1	CS-101
2	CS-347
3	PHY-101

UNION

	course_id
1	CS-101
2	CS-315
3	CS-319
4	CS-319
5	FIN-201
6	HIS-351
7	MU-199



	course_id
1	CS-101
2	CS-315
3	CS-319
4	CS-347
5	FIN-201
6	HIS-351
7	MU-199
8	PHY-101



## Set operations in SQL – Union All

- Use **UNION ALL** instead of **UNION** to retain all duplicates

```
(select course_id
  from section
  where semester = 'Fall' and year = 2017)
union all
(select course_id
  from section
  where semester = 'Spring' and year = 2018)
```

	course_id
1	CS-101
2	CS-347
3	PHY-101

UNION ALL



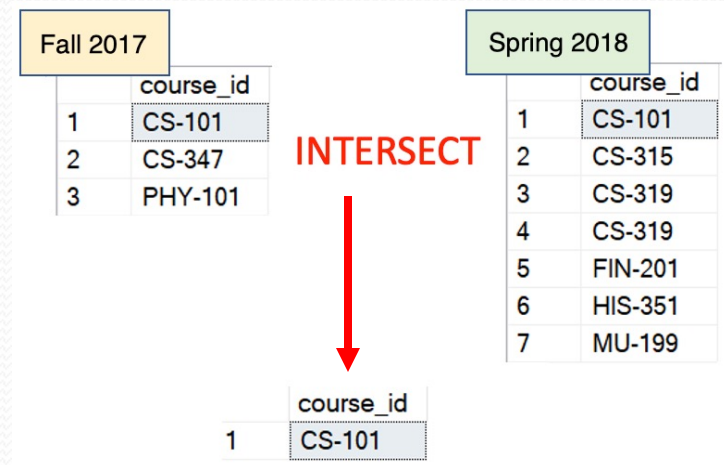
	course_id
1	CS-101
2	CS-347
3	PHY-101
4	CS-101
5	CS-315
6	CS-319
7	CS-319
8	FIN-201
9	HIS-351
10	MU-199

	course_id
1	CS-101
2	CS-315
3	CS-319
4	CS-319
5	FIN-201
6	HIS-351
7	MU-199

## Set operations in SQL – Intersect

- **INTERSECT** of two tables A and B, is a table containing all rows that are common to A and B
- E.g., list all the courses that are taught in both Fall 2017 and Spring 2018

```
(select course_id
  from section
  where semester = 'Fall' and year = 2017)
intersect
(select course_id
  from section
  where semester = 'Spring' and year = 2018)
```





## Set operations in SQL – Intersect All

- The **INTERSECT** operation automatically eliminates duplicates
- Use **INTERSECT ALL** instead to keep all duplicates

```
(select course_id
  from section
 where semester = 'Fall' and year = 2017)
intersect all -- SQL Server does not support
(select course_id
  from section
 where semester = 'Spring' and year = 2018)
```

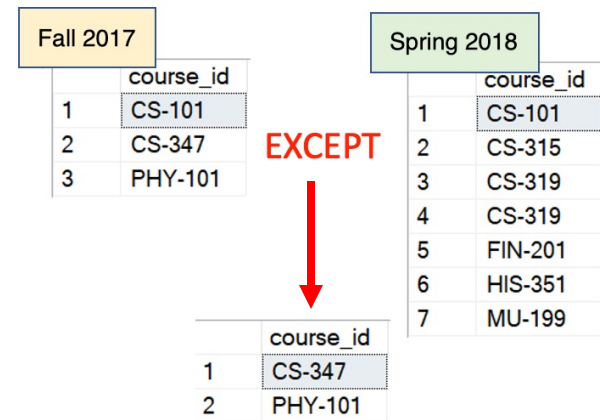
1	CS-101	01	Fall	2017
2	CS-101	02	Fall	2017
3	SE-340	01	Fall	2017
4	SE-340	02	Fall	2017
5	SE-340	03	Fall	2017
6	SE-340	04	Fall	2017
7	CS-101	01	Spring	2018
8	CS-101	02	Spring	2018
9	SE-340	01	Spring	2018
10	SE-340	02	Spring	2018

	course_id
	character varying (10)
1	SE-340
2	SE-340
3	CS-101
4	CS-101

## Set operations in SQL – Except

- **EXCEPT** of two tables A and B, is a table containing all rows that are in A but not in B
- E.g., find all courses taught in Fall 2017 but not in Spring 2018

```
(select course_id
  from section
  where semester = 'Fall' and year = 2017)
except
(select course_id
  from section
  where semester = 'Spring' and year = 2018)
```





## Set operations in SQL – Except All

- **EXCEPT** operation automatically eliminates duplicates
- Use **EXCEPT ALL** instead to retain duplicates

```
(select course_id
  from section
  where semester = 'Fall' and year = 2017)
except all -- SQL Server does not support
(select course_id
  from section
  where semester = 'Spring' and year = 2018)
```

1	CS-101	01	Fall	2017
2	CS-101	02	Fall	2017
3	SE-340	01	Fall	2017
4	SE-340	02	Fall	2017
5	SE-340	03	Fall	2017
6	SE-340	04	Fall	2017
7	CS-101	01	Spring	2018
8	CS-101	02	Spring	2018
9	SE-340	01	Spring	2018
10	SE-340	02	Spring	2018

	course_id
	character varying ( )
1	SE-340
2	SE-340

# Modification of the database

- **Deletion**
- **Insertion**
- **Updates**



# Deletion

- The **DELETE** statement allows rows to be deleted from a named table
- The format of **DELETE**:

```
delete from tableName          -- only one relation  
    [where searchCondition]    -- optional
```

- **WHERE** clause can be as complex as a **SELECT** command's **WHERE** clause

# Deletion

- E.g., delete all tuples from the instructor relation

*delete from instructor;*

- All tuples are deleted
- The instructor relation still exists
  - Empty



# Deletion

- E.g., delete all instructor's information from Finance department

```
delete from instructor  
  where dept_name = 'Finance';
```

- E.g., delete all instructors with a salary between \$70,000 and \$90,000

```
delete from instructor  
  where salary between 70000 and 90000;
```

# Deletion

- E.g., delete all instructors whose dept located in the Watson building

```
delete from instructor
  where dept_name in (select dept_name
                     from department
                     where building = 'Watson');
```

- First, find all departments located in Watson
- Then, delete tuples pertaining to those departments



## Deletion

- E.g., delete all the records of all instructors with salary below the average at the university

```
delete from instructor
  where salary < (select AVG(salary)
                 from instructor);
```

- First, test each tuple in the relation
- Then, delete all those tuples passing the test
  - instructors with a salary lower than average

# Insertion

- To insert data into a relation
  - Specify a tuple to be inserted
    - **INSERT ... VALUES**
  - Write a query whose result is a set of tuples to be inserted
    - **INSERT ... SELECT**



## Insertion – specify tuple to be inserted

- E.g., insert a course CS-484 in the Computer Science department

```
insert into course  
  values ( 'CS-484', 'Database Management', 'Comp. Sci.', 3 );
```

- The values are specified in the order in which the corresponding attributes are listed in the relation schema

	course_id	title	dept_name	credits
5	CS-190	Game Design	Comp. Sci.	4
6	CS-315	Robotics	Comp. Sci.	3
7	CS-319	Image Processing	Comp. Sci.	3
8	CS-347	Database System Concepts	Comp. Sci.	3
9	CS-484	Database Management	Comp. Sci.	3
10	EE-181	Intro. to Digital Svstems	Elec. Ena.	3

## Insertion – specify tuple to be inserted

- SQL allows the attributes to be specified as part of the **INSERT** statement

```
insert into course ( course_id, title, dept_name, credits )  
  values ( 'CS-484', 'Database Management', 'Comp. Sci.', 3 );
```

```
-- same
```

```
insert into course ( title, course_id, credits, dept_name )  
  values ( 'Database Management', 'CS-484', 3, 'Comp. Sci.' );
```



## Insertion – specify tuple to be inserted

- If values are given on only some attributes, the remaining attributes are assigned a null value

```
insert into student  
values ( '3003', 'Green', 'Finance', null );
```

## Insert – Select

- Insert tuples based on the result of a query
- E.g., for those students in music dept who earned more than 144 credits, make them to be instructors in the music dept with a salary of \$18,000

```
insert into instructor
  select ID, name, dept_name, 18000  -- constants are allowed
  from student                      -- value applied to all tuples
  where dept_name = 'Music' and tot_cred > 144;
```

- The system evaluates the **SELECT** statement fully before it performs any insertions
- The resulting relation must be consistent with the instructor schema



# Insertion

- Most RDBMS have special “bulk loader” utilities to insert a large set of tuples into a relation
- E.g., read a CSV file into a relation in SQL Server

```
BULK INSERT section
FROM '/Users/ken/Desktop/484/import.csv'      -- macOS path
WITH
(
    FIRSTROW          = 1,
    FIELDTERMINATOR    = ',', -- CSV field delimiter
    ROWTERMINATOR      = '\n', -- shifts control to the next row
    TABLOCK
)
```

# Update

- The **UPDATE** statement allows the contents of existing rows in a named table to be changed.

```
update tableName
  set columnName1 = dataValue1 [, columnName2 = dataValue2....]
  [where searchCondition] -- optional
                        -- if omitted, all named columns for all
                        -- tuples are updated
                        -- if specified, only those tuples that satisfy
                        -- the searchCondition are updated
```

- Specify the names of one or more columns to be updated



# Update

- **WHERE** clause
  - In general, the **WHERE** clause of the **UPDATE** statement may contain any legal construct legal found in the **WHERE** clause of the **SELECT** statement

# Update

- E.g., increase all instructor's salary by 5 percent

```
update instructor
  set salary = salary * 1.05;
-- WHERE clause is omitted
-- all tuples will be updated
```



# Update

- E.g., apply a 5% salary increase to those instructors whose salaries are less than \$70,000

```
update instructor
  set salary = salary * 1.05;
  where salary < 70000;    -- tuples with salary < 70000 are updated
```

# Update

- E.g., apply a 5% salary increase to those instructors whose salaries are less than the average instructor salary

```
update instructor
  set salary = salary * 1.05;
where salary < (select AVG(salary)
from instructor);      -- tuples with salary < 72000 are updated
```



# Update

- E.g., apply a salary increase to all instructors
  - 3% for salaries greater than \$100,000
  - 5% for all others

```
update instructor
  set salary = salary * 1.03;
  where salary > 100000;
```

-- the order of these two statements is important

```
update instructor
  set salary = salary * 1.05;
  where salary <= 100000;
```

# Update

- SQL provides a **CASE** construct that we can use to perform both updates with a single **UPDATE** statement, avoiding the problem with the order of updates

```
CASE                                -- can be used anywhere
  WHEN pred1 THEN result1          -- a value is expected
  WHEN pred2 THEN result2
  ...
  WHEN predn THEN resultn
  ELSE result0
END
```



# Update

- E.g., apply a salary increase to all instructors
  - 3% for salaries greater than \$100,000
  - 5% for all others

```
update instructor
  set salary = case
    when salary > 100000 then salary * 1.03
    else salary * 1.05
  end;
```

# Acknowledgements

- WIKIPEDIA
  - <https://en.wikipedia.org/wiki/SQL>