

CSc 484

Database Management Systems

Ken Gamradt

Spring 2024

Database Design using ER (II)

Design Phases

- **Initial phase:** characterize fully the data needs of the prospective database users
- **Second phase:** choosing a data model
 - Applying the concepts of the chosen data model
 - Translating these requirements into a conceptual schema of the database
 - A fully developed conceptual schema indicates the functional requirements of the enterprise
 - Describe the kinds of operations (or transactions) that will be performed on the data

Design Phases

- **Final Phase:** Moving from an abstract data model to the implementation of the database
 - **Logical Design:** deciding on the database schema
 - Database design requires that we find a “good” collection of relation schemas
 - Business decision: What attributes should we record in the database
 - Computer Science decision:
 - What relation schemas should we have
 - How should the attributes be distributed among the various relation schemas
 - **Physical Design:** deciding on the physical layout of the database

Design Alternatives

- In designing a database schema, we must ensure that we avoid two major pitfalls:
 - **Redundancy:** a bad design may result in repeated information
 - Redundant representation of information may lead to data inconsistency among the various copies of information
 - **Incompleteness:** a bad design may make certain aspects of the enterprise difficult or impossible to model
- Avoiding bad designs is not enough
 - There may be many good designs from which we must choose

Design Approaches

- **Entity Relationship Model (Chapter 6)**
 - Models an enterprise as a collection of **entities** and **relationships**
 - **Entity**: a “**thing**” or “**object**” in the enterprise that is **distinguishable** from other objects
 - Described by a set of attributes
 - **Relationship**: an association among several entities
 - Represented diagrammatically by an entity-relationship diagram:
- **Normalization Theory (Chapter 7)**
 - Formalize what designs are bad, and test for them

Entity-Relationship Model



- The ER data model was developed to facilitate database design by allowing specification of an enterprise schema that represents the overall logical structure of a database
- The ER data model employs three basic concepts:
 - **Entity sets**
 - **Relationship sets**
 - **Attributes**
- The ER model also has an associated diagrammatic representation, the ER diagram, which can express the overall logical structure of a database graphically

Entity Sets

- An **entity** is an object that exists and is distinguishable from other objects
 - E.g., specific person, company, event, plant, ...
- An **entity set** is a set of entities of the same type that share the same properties
 - E.g., set of all persons, companies, trees, holidays, ...
- An entity is represented by a set of attributes; i.e., descriptive properties possessed by all members of an entity set
 - E.g.,
instructor = (ID, name, salary)
course = (course_id, title, credits)
- A subset of the attributes form a **primary key** of the entity set
 - I.e., uniquely identifying each member of the set

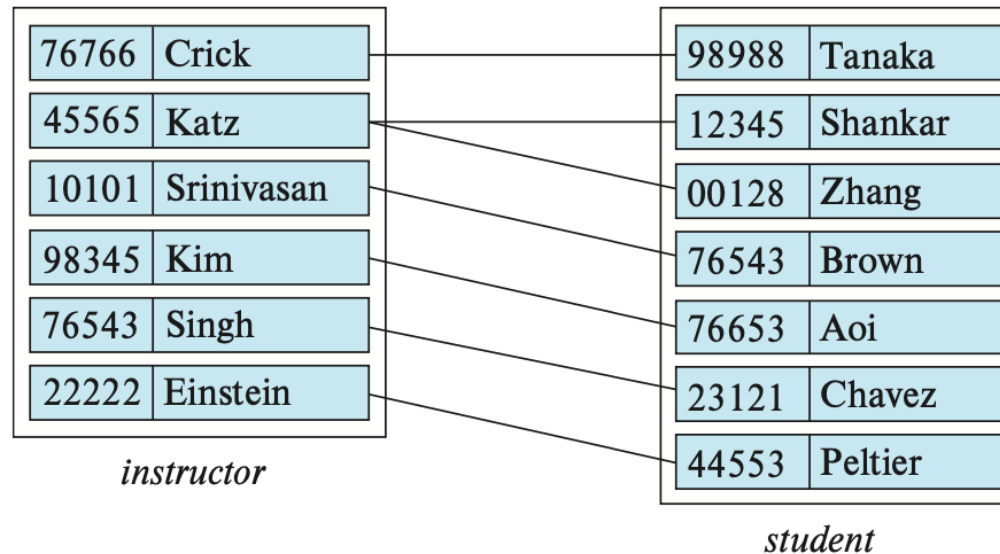
Representing Entity sets in E-R Diagram

- **Entity set**
 - Set of entities of the same type that share the same attributes
- An entity set is represented in an E-R diagram by a rectangle
 - Part 1: name of the entity set
 - Instructor
 - Student
 - Part 2: name of all **attributes**
 - ID, name, salary
 - ID, name, tot_cred
 - Primary keys are underlined
 - ID
 - ID




Relationship Sets

- E.g., we define the relationship set *advisor* to denote the associations between students and the instructors who act as their advisors
- Pictorially, we draw a line between related entities



Relationship Sets

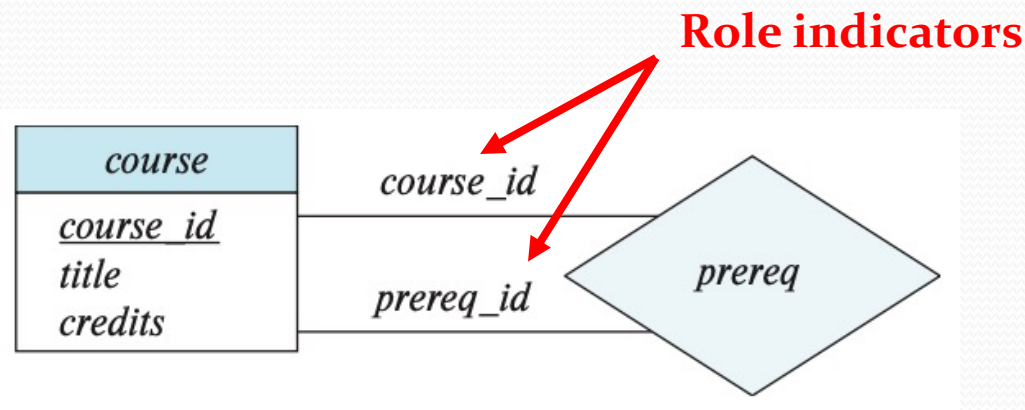
- **Relationship set**
 - Is represented in an E-R diagram by a diamond
 - Link the entity sets via lines



- Each entity plays a role in the relationship

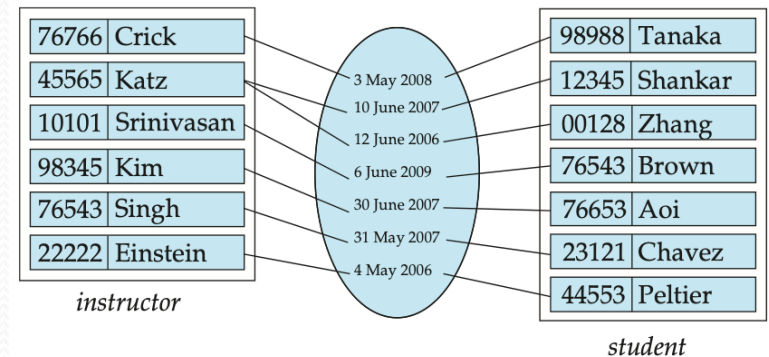

Recursive Relationship Set

- The same entity set participates in a relationship set more than once
 - In different roles
- E.g., depict the prerequisite course for courses



Relationship Sets with Attributes

- A relationship may have attributes – **descriptive attributes**
- Attributes of a relationship set are represented using an **undivided** rectangle
- Linked by dashed line from relationship set to descriptive attribute




Degree of a Relationship Set

- **Binary** relationship
 - involve two entity sets
 - Degree of two
 - most relationship sets in a database system are binary
- Relationships between more than two entity sets are rare
 - E.g., students work on research projects under the guidance of an instructor
 - relationship *proj_guide* is a **ternary** relationship between instructor, student, and project

Non-binary Relationship Sets

- Most relationship sets are binary
- There are occasions when it is more convenient to represent relationships as non-binary
- E-R Diagram with a **Ternary** Relationship

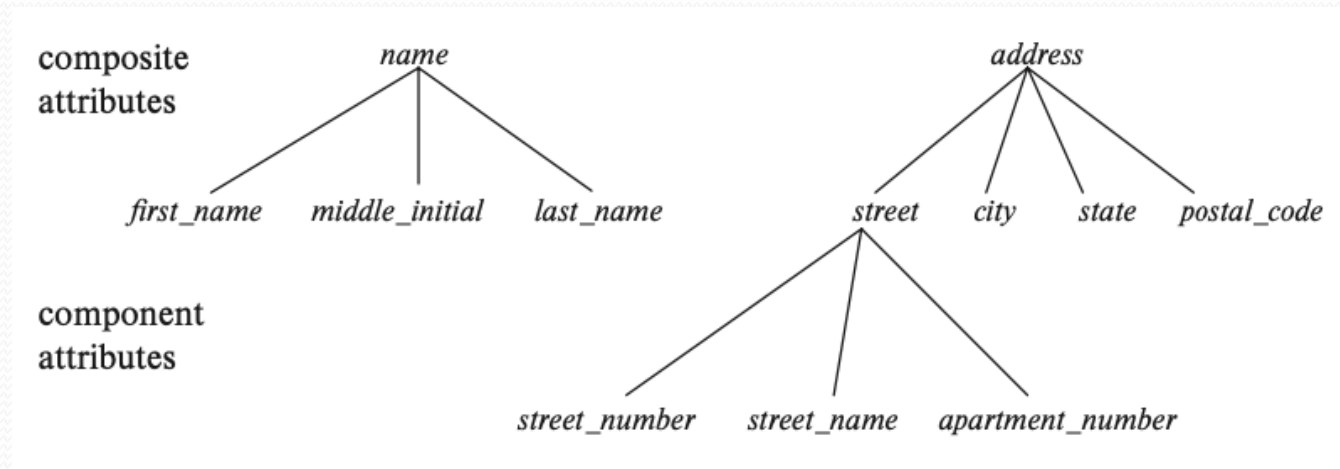


Complex Attributes

- Attribute types:
 - **Simple** and **composite** attributes
 - **Single-valued** and **multivalued** attributes
 - E.g., multivalued attribute: *phone_numbers*
 - **Derived** attributes
 - Can be computed from other attributes
 - E.g., *age* – given *date_of_birth*
- **Domain** – the set of permitted values for each attribute

Composite Attributes

- **Simple** attribute:
 - Cannot be decomposed
 - Broken into smaller parts
- **Composite** attribute:
 - Can be decomposed into other meaningful attributes



Representing Complex Attributes

- If the name and address attributes are added to the instructor entity-set
- **When to use composite attributes?**
- If the users wish to refer to
 - an entire attribute on some occasions
 - a component of other occasions
- Another benefit of using composite attributes
 - Help to group together related attributes, making the modeling cleaner

<i>instructor</i>
<u><i>ID</i></u>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
<i>city</i>
<i>state</i>
<i>zip</i>
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age</i> ()

Representing Complex Attributes

- **Single-valued** attribute:
 - Has exactly one value for a particular entity
- **Multi-valued** attribute:
 - May have multiple values for a particular entity
 - E.g., an instructor may have zero or more phone numbers, and different instructors may have different numbers of phone numbers
- **Derived** attribute:
 - The value of this type of attribute can be derived from the values of other related attributes or entities

<i>instructor</i>
<u><i>ID</i></u>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
<i>city</i>
<i>state</i>
<i>zip</i>
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age</i> ()

Structural Constraints


- The constraints that may be placed on entity types that participate in a relationship
 - Reflect the restrictions on the relationships as perceived in the “real world”
- Two types of constraint on relationship sets
 - **Cardinality**
 - **Participation**

Mapping Cardinalities

- **Mapping cardinalities** (cardinality ratios)
 - Express the number of entities to which another entity can be associated via a relationship set
 - Specify the minimum and/or maximum number of each entities when participating a relationship
 - Minimum cardinality: 0, 1
 - Maximum cardinality: 1, N

Entity-Relationship Model – Mapping Cardinalities



- Describe a binary relationship set
- E.g., **advisor** relationship set between **instructor** and **student** entity sets
- The mapping cardinality can be
 - **One-to-one**
 - **One-to-many**
 - **Many-to-one**
 - **Many-to-many**
- A directed line and an undirected line are used to indicate the cardinality constraint in the E-R diagram



Mapping Cardinalities

- **One-to-one**

- Each instructor may advise **at most one** student
- Each student may have **at most one** advisor





ER diagram notation:

Directed line from the relationship set to the “**one**” side of the relationship

Mapping Cardinalities

- **One-to-many**

- Each instructor entity is associated with any number (**zero or more**) of student entities
- Each student entity is associated with **at most one** instructor entity





ER diagram notation:

Undirected line from the relationship set to the “**many**” side of the relationship

Mapping Cardinalities

- **Many-to-one**



- Each instructor may advise **at most one** student
- Each student may have any number (**zero or more**) of advisors



Mapping Cardinalities

- **Many-to-many**

- Each instructor may advise any number (**zero or more**) of students
- Each student may have any number (**zero or more**) of advisors




Participation

- Determine whether all or only some entities of an entity set participate in a relationship
 - **Total** participation
 - Every entity in entity set must participate in at least one relationship
 - **Partial** participation
 - Some entities do not have to participate in relationship
- **Double lines** are used in E-R diagrams to indicate **total** participation

Participation

- E.g., the university requires that each student must have an advisor




Cardinality & Participation

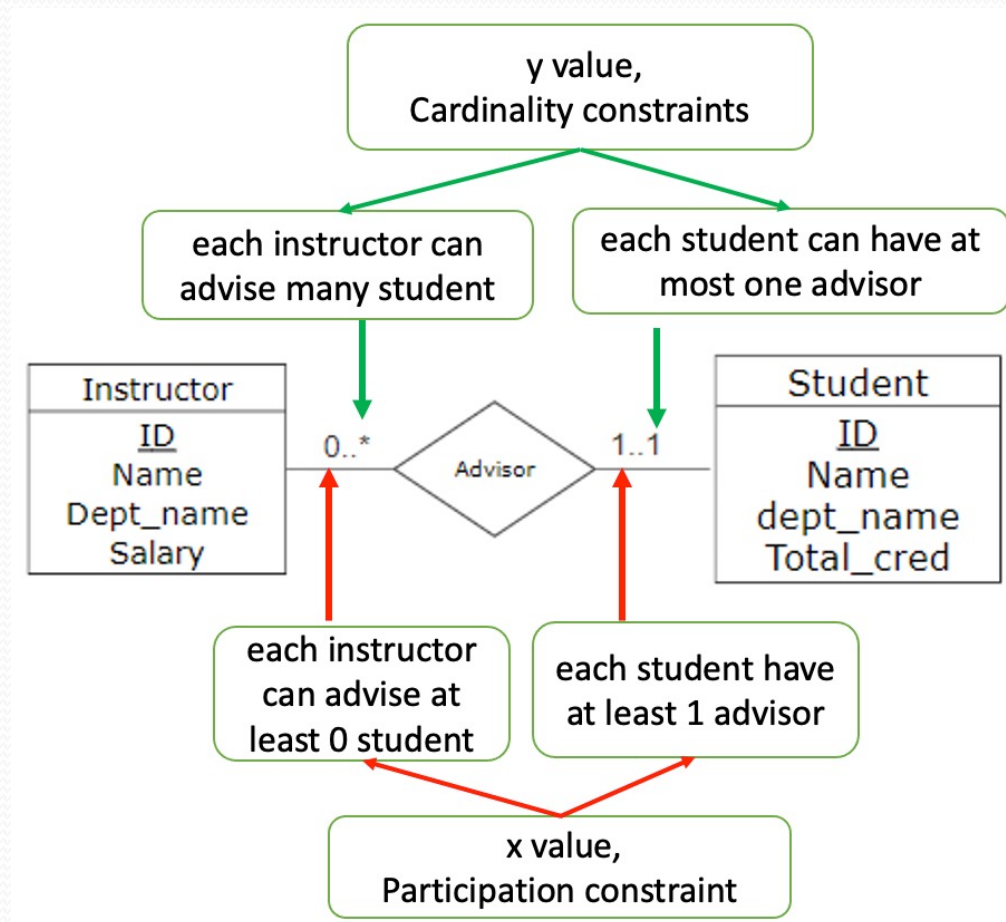
- Another way to indicate more complex constraints of cardinality and participation
 - Single line together with numbers (x .. y) to indicate the constraints from the requirement
 - x: meaning minimum cardinality
 - y: meaning maximum cardinality
 - x, y can be numbers (0, 1, 2, ...) and * (means unlimited)
 - x = 1: total participation
 - x = 0: partial participation
 - y = 1: the entity participates in at most one relationship
 - y = *: unlimited number of entities

Cardinality & Participation

- The university requires
 - Each student must have at exactly one advisor
 - Each instructor can advise any number of students
- Cardinality: one-to-many from instructor to student
 - For instructors that participate
- Participation: instructor (**partial**), student (**total**)




Cardinality & Participation



Cardinality & Participation

- E.g., university requires
 - Each instructor must advise at least one student
 - Each student must have exactly one advisor
- Both entity sets are total participation
 - Cardinality is one-to-many





Cardinality Constraints on Ternary Relationship

- We allow at most one arrow out of a ternary (or greater degree) relationship to indicate a cardinality constraint
- E.g., an arrow from proj_guide to instructor indicates each student has at most one guide for a project
- If there is more than one arrow, there are two ways of defining the meaning
 - E.g., a ternary relationship R between A, B and C with arrows to B and C could mean
 1. Each A entity is associated with a unique entity from B and C or
 2. Each pair of entities from (A, B) is associated with a unique C entity, and each pair (A, C) is associated with a unique B
 - Each alternative has been used in different formalisms
 - To avoid confusion, we outlaw more than one arrow

Cardinality Constraints on Ternary Relationship

- Constraints on non-binary relationship sets
- E.g., senior design projects
 - Instructors provide guidance for any number of students on their projects
 - Each student can be guided by exactly one instructor

Only one arrow out of a non-binary relationship set is allowed



Primary key

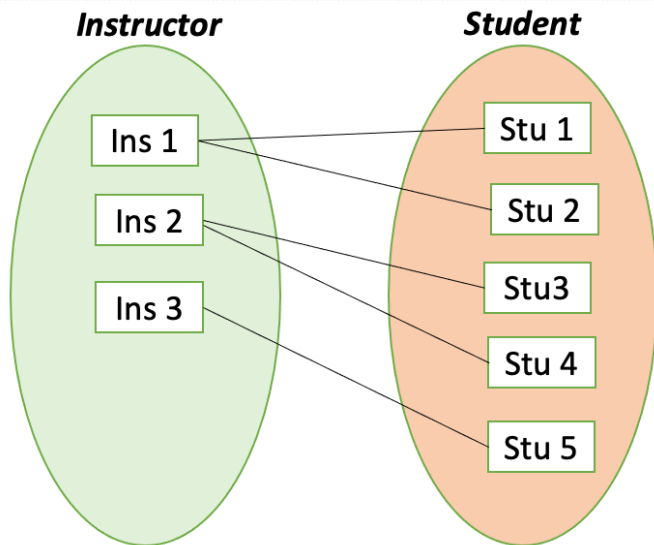
- Primary keys provide a way to specify how entities and relations are distinguished
 - **Entity sets**
 - **Relationship sets**
 - **Weak entity sets**

Primary key for Entity Sets

- Conceptually, individual entities are distinct
- From database perspective, the differences among them must be expressed in terms of their attributes
- The values of the attribute values of an entity must be such that they can uniquely identify the entity
 - No two entities in an entity set are allowed to have exactly the same value for all attributes
- A key for an entity is a set of attributes that can be used to distinguish entities from each other

Primary key for Relationship sets


- Keys also help to identify relationships uniquely



Each relationship is unique

Entity set names used as a prefix to distinguish attributes with the same name, can be omitted if attributes with different names


(Instructor.ID, Student.ID) form a superkey for the relationship key



primary-key(E_1) U primary-key(E_2) U ... U primary-key(E_n) can be used as a superkey to distinguish relationships

Primary key for Relationship sets

- More about the name of attributes in the superkey of a relationship set
 - In recursive relationship set, one entity set plays different roles in the relationship
- Use role names/indicators as prefixes to form a unique attribute



(*course_id.ID*, *prereq_id.ID*) forms the superkey for *prereq*


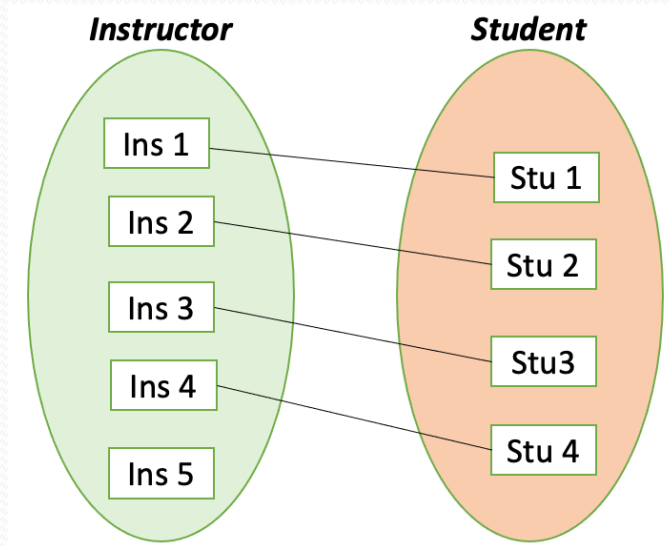
Primary key for Relationship sets

- The choice of the primary key depends on the mapping cardinality of the relationship
- Considering a binary relationship set

Primary key for Relationship sets

- **One-to-one**

- The primary key of the relationship set can be the primary key of either one of the participating entity sets


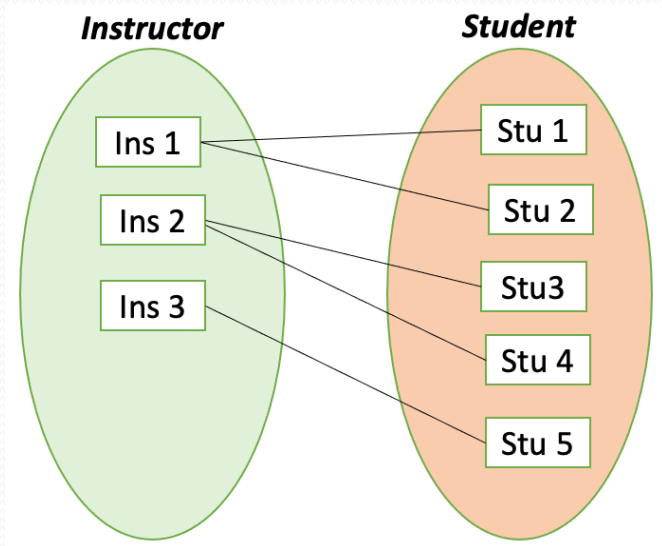


The primary key for advisor is:

Instructor.ID or ***Student.ID***

Primary key for Relationship sets


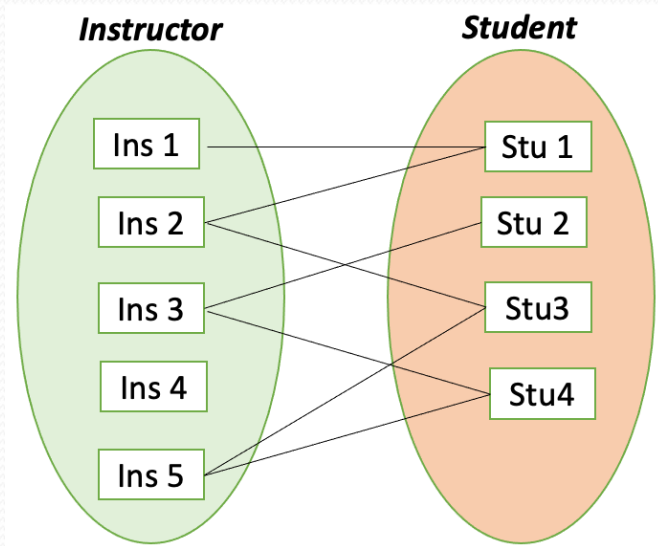
- **One-to-many**
 - The primary key for the “many” side is used as the primary key for the relationship set



The primary key for advisor is:
Student.ID

Primary key for Relationship sets

- **Many-to-many**
 - The primary key is the union of both entity sets



The primary key for advisor is:
(Instructor.ID, Student.ID)



Acknowledgements