

CSc 484

# Database Management Systems

Ken Gamradt

Spring 2024

Accessing SQL from a Programming Language

# Accessing SQL from a Programming Language

- A database programmer must have access to a general-purpose programming language for at least two reasons
  - Not all queries can be expressed in SQL
    - SQL does not provide the full expressive power of a general-purpose language
  - Non-declarative actions such as
    - printing a report
    - interacting with a user
    - sending the results of a query to a graphical user interfacecannot be done from within SQL



# Accessing SQL from a Programming Language

- There are two approaches to accessing SQL from a general-purpose programming language
  - A general-purpose program
    - Can connect to and communicate with a database server using a collection of functions
  - Embedded SQL
    - Provides a means by which a program can interact with a database server
      - SQL statements are translated at compile time into function calls
      - At runtime, these function calls connect to the database using an API that provides dynamic SQL facilities

# Accessing SQL from a Programming Language

- Dynamic SQL:
  - Allows the program to construct an SQL query as a character string at runtime, submit the query, and then retrieve the result into program variables a tuple at a time
  - String query = “SELECT \* FROM instructor;”
- Embedded SQL:
  - The SQL statements are identified at compile time using a preprocessor, which translates requests expressed in embedded SQL into function calls
  - At runtime, the function calls connect to the database
  - EXEC SQL SELECT \* FROM instructor;



# JDBC

- **JDBC** is a Java API for communicating with database systems supporting SQL
  - **Java Database Connectivity**
- JDBC supports a variety of features for
  - Querying data
  - Updating data
  - Retrieving query results
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes.
- Model for communicating with the database:
  - Open a connection
  - Create a “statement” object
  - Execute queries using the statement object to send queries and fetch results
  - Exception mechanism to handle errors

# JDBC

```
import java.sql.*;
// import java.sql.Connection;
// import java.sql.DriverManager;
// import java.sql.SQLException;
// import java.sql.ResultSet;
// import java.sql.Statement;
// import java.sql.PreparedStatement;

public static void JDBCexample(
    String userid, String passwd) {
    try {
        // Class.forName("org.postgresql.Driver");
        Connection conn = DriverManager.getConnection(
            "jdbc:postgresql://host:5432/university",
            // "jdbc:mysql://host:3306/database",
            // "jdbc:sqlite:database_file_path",
            // "jdbc:sqlserver://host:53000;databaseName=database",
            userid,
            passwd);
        Statement stmt = conn.createStatement();

        // Do something

    } catch (Exception sqle) {
        System.out.println("Exception: " + sqle);
    } finally {
        if (rset != null) try { rset.close(); } catch (SQLException ignore) {}
        if (stmt != null) try { stmt.close(); } catch (SQLException ignore) {}
        if (conn != null) try { conn.close(); } catch (SQLException ignore) {}
    }
}
```



# JDBC

```
// Update to database
try {
    stmt.executeUpdate(
        "insert into instructor values(" +
        "'77987', 'Kim', 'Physics', 98000)");
}
catch (SQLException sqle) {
    System.out.println(
        "Could not insert tuple " + sqle);
}
// Execute query and fetch and print results
ResultSet rset = stmt.executeQuery(
    "select dept_name, avg(salary) " +
    "from instructor " +
    "group by dept");
while (rset.next()) {
    System.out.println(
        rset.getString("dept_name") + " " +
        rset.getFloat(2));
}
```

# JDBC Subsections

- Connecting to the Database
- Shipping SQL Statements to the Database System
- Exceptions and Resource Management
- Retrieving the Result of a Query
- Prepared Statements
- Callable Statements
- Metadata Features
- Other Features
- Database Access from Python



## JDBC Code Details

- Getting result fields:

```
// equivalent if dept_name is the first argument of select result  
rset.getString("dept_name") and rset.getString(1)
```

- Dealing with Null values

```
int a = rset.getInt("a");  
if (rset.isNull())  
    Systems.out.println("Got null value");
```

## Prepared Statements

```
PreparedStatement pstmt = conn.prepareStatement(
    "insert into instructor values(?,?,?,?)");
pstmt.setString(1, "88877");
pstmt.setString(2, "Perry");
pstmt.setString(3, "Finance");
pstmt.setInt(4, 125000);
pstmt.executeUpdate();
pstmt.setString(1, "88878");
pstmt.executeUpdate();
```

- WARNING: always use prepared statements when taking an input from the user and adding it to a query
  - // NEVER create a query by concatenating strings
  - "insert into instructor values(' " + ID + " ', ' " + name + " ', " + " ' + dept name + " ', " ' balance + ')"
  - What if name is "D'Souza"?



# SQL Injection

- Suppose query is constructed using
  - "select \* from instructor where name = " + name + ""
- Suppose the user, instead of entering a name, enters:
  - X' or 'Y' = 'Y
- then the resulting statement becomes:
  - "select \* from instructor where name = " + "X' or 'Y' = 'Y" + ""
  - which is:
    - select \* from instructor where name = 'X' or 'Y' = 'Y'
  - User could have even used
    - X'; update instructor set salary = salary + 10000; --
- Prepared statement internally uses:
  - "select \* from instructor where name = 'X\' or \'Y\' = \'Y'
- **Note:** Always use prepared statements, with user inputs as parameters

# Metadata

- ResultSet metadata
- E.g., after executing query to get a ResultSet rset:

```
ResultSetMetaData rsmd = rset.getMetaData();  
for(int i = 1; i <= rsmd.getColumnCount(); i++) {  
    System.out.println(rsmd.getColumnName(i));  
    System.out.println(rsmd.getColumnTypeName(i));  
}
```



# Metadata

- Database metadata

```
DatabaseMetaData dbmd = conn.getMetaData();
// Arguments to getColumns: Catalog, Schema-pattern, Table-pattern,
// and Column-Pattern
// Returns: One row for each column; row has a number of attributes
// such as COLUMN_NAME, TYPE_NAME
// The value null indicates all Catalogs/Schemas
// The value "" indicates current catalog/schema
// The value "%" has the same meaning as SQL like clause
ResultSet rs = dbmd.getColumns(null, "university", "department", "%");
while( rs.next()) {
    System.out.println(rs.getString("COLUMN_NAME"), rs.getString("TYPE_NAME"));
}
```

# Metadata

- Database metadata

```
DatabaseMetaData dbmd = conn.getMetaData();
// Arguments to getTables: Catalog, Schema-pattern, Table-pattern,
// and Table-Type
// Returns: One row for each table; row has a number of attributes
// such as TABLE_NAME, TABLE_CAT, TABLE_TYPE, ...
// The value null indicates all Catalogs/Schemas
// The value "" indicates current catalog/schema
// The value "%" has the same meaning as SQL like clause
// The last attribute is an array of types of tables to return
// TABLE means only regular tables
ResultSet rs = dbmd.getTables("", "", "%", new String[] {"TABLES"});
while( rs.next()) {
    System.out.println(rs.getString("TABLE_NAME"));
}
```



## Finding Primary Keys

```
DatabaseMetaData dmd = connection.getMetaData();
// Arguments below are: Catalog, Schema, and Table
// The value "" for Catalog/Schema indicates current catalog/schema
// The value null indicates all catalogs/schemas
ResultSet rs = dmd.getPrimaryKeys("", "", tableName);
while(rs.next()) {
    // KEY_SEQ indicates the position of the attribute in the
    // primary key, which is required if a primary key has multiple
    // attributes
    System.out.println(rs.getString("KEY_SEQ"), rs.getString("COLUMN_NAME"));
}
```

## Other JDBC Features

- Calling functions and procedures
  - CallableStatement cStmt1 =  
    conn.prepareCall("{? = call some function(?)})");
  - CallableStatement cStmt2 =  
    conn.prepareCall("{call some procedure(?,?)})");
- Handling large object types
  - getBlob() and getClob() that are similar to the getString() method, but return objects of type Blob and Clob, respectively
  - get data from these objects by getBytes()
  - associate an open stream with Java Blob or Clob object to update large objects
    - blob.setBlob(int parameterIndex, InputStream inputStream)



## Other JDBC Features

- A BLOB is binary large object that can hold a variable amount of data with a maximum length of 65535 characters
- These are used to store large amounts of binary data, such as
  - images
  - other types of files
- Fields defined as TEXT also hold large amounts of data
- Difference between the two is the sorts and comparisons done on the stored data
  - case sensitive on BLOBs
  - not case sensitive in TEXT fields
- You do not specify a length with BLOB or TEXT

## Other JDBC Features

- CLOB stands for Character Large Object
- SQL Clob is a built-in datatype and is used to store large amount of textual data
  - Can store data up to 2,147,483,647 characters
- The **java.sql.Clob** interface of the JDBC API represents the CLOB datatype
- Since the Clob object in JDBC is implemented using an SQL locator, it holds a logical pointer to the SQL CLOB (not the data)
- MySQL database provides support for this datatype using four variables
  - **TINYTEXT:** maximum of  $2^8-1$  (255) characters
  - **TEXT:** maximum of  $2^{16}-1$  (65535) characters
  - **MEDIUMTEXT:** maximum of  $2^{24}-1$  (16777215) characters
  - **LONGTEXT:** maximum of  $2^{32}-1$  (4294967295) characters



## Transaction Control in JDBC

- By default, each SQL statement is treated as a separate transaction that is committed automatically
  - Bad idea for transactions with multiple updates
- Can turn off automatic commit on a connection
  - `conn.setAutoCommit(false);`
- Transactions must then be committed or rolled back explicitly
  - `conn.commit();` or
  - `conn.rollback();`
- `conn.setAutoCommit(true)` turns on automatic commit

# Transaction Control in JDBC

```
import java.sql.*;

public class BatchProcessing_Statement {
    public static void main(String args[]) throws Exception {
        // Getting the connection
        String mysqlUrl = "jdbc:mysql://localhost/sampleDB";
        Connection conn = DriverManager.getConnection(mysqlUrl, "root", "password");
        System.out.println("Connection established.....");
        // CREATE TABLE Dispatches( Product_Name VARCHAR(255), Name_Of_Customer
        VARCHAR(255), Month_Of_Dispatch VARCHAR(255), Price INT, Location VARCHAR(255));
        // Creating a Statement object
        Statement stmt = conn.createStatement();
        // Setting auto-commit false
        conn.setAutoCommit(false);
        // Statements to insert records
        String insert1 = "INSERT INTO Dispatches( Product_Name , Name_Of_Customer , "
            + "Month_Of_Dispatch, Price, Location) VALUES "
            + "('Keyboard', 'Amith', 'January', 1000, 'hyderabad')";
        String insert2 = "INSERT INTO Dispatches( Product_Name , Name_Of_Customer , "
            + "Month_Of_Dispatch , Price, Location) VALUES "
            + "('Mouse', 'Sudha', 'September', 200, 'Vijayawada')";
        // Adding the statements to the batch
        stmt.addBatch(insert1);
        stmt.addBatch(insert2);
        // Executing the batch
        stmt.executeBatch();
        // Saving the changes
        conn.commit();
        System.out.println("Records inserted.....");
    }
}
```



# Database Access from Python

```
import psycopg2 # PostgreSQL database adapter for Python

def PythonDatabaseExample(userid, passwd)
    try:
        conn = psycopg2.connect(
            host="host",
            port=5432,
            dbname="university",
            user=userid,
            password=passwd)
        cur = conn.cursor()
        try:
            cur.execute("insert into instructor values(%s, %s, %s, %s)",
                        ("77987", "Kim", "Physics", 98000))
        conn.commit()
    except Exception as sqle:
        print("Could not insert tuple. ", sqle)
        conn.rollback()
    cur.execute(("select dept_name, avg(salary) "
                "from instructor group by dept_name"))
    for dept in cur:
        print dept[0], dept[1]
    except Exception as sqle:
        print("Exception : ", sqle)
```

## Database Access from Python – The cursor class

- Allows Python code to execute PostgreSQL commands in a database session
- Cursors are created by the `connection.cursor()` method:
  - they are bound to the connection for the entire lifetime
  - all the commands are executed in the context of the database session wrapped by the connection
- Cursors created from the same connection are not isolated
  - any changes done to the database by a cursor are immediately visible by the other cursors
- Cursors created from different connections can or can not be isolated, depending on the connections' `isolation level`
  - See also `rollback()` and `commit()` methods



# ODBC

- **Open DataBase Connectivity** (ODBC) standard
  - standard for application programs to communicate with a database server
  - application program interface (API) to
    - open a connection with a database
    - send queries and updates
    - get back results
- Applications such as GUI, spreadsheets, ... can use ODBC

# ODBC

```
void ODBCexample() {
    RETCODE error;
    HENV env;    /* environment */
    HDBC conn;   /* database connection */
    SQLAllocEnv(&env);
    SQLAllocConnect(env, &conn);    /* SQL_NTS - null-terminated string */
    SQLConnect(conn, "host", SQL_NTS, "avi", SQL_NTS, "avipasswd", SQL_NTS);
    {
        char deptname[80];
        float salary;
        int lenOut1, lenOut2;
        HSTMT stmt;
        char * sqlquery =
            "select dept name, sum (salary) from instructor group by dept name";
        SQLAllocStmt(conn, &stmt);
        error = SQLExecDirect(stmt, sqlquery, SQL_NTS);
        if (error == SQL_SUCCESS) {
            SQLBindCol(stmt, 1, SQL_C_CHAR, deptname, 80, &lenOut1);
            SQLBindCol(stmt, 2, SQL_C_FLOAT, &salary, 0, &lenOut2);
            while (SQLFetch(stmt) == SQL_SUCCESS) {
                printf (" %s %g\n", deptname, salary);
            }
            SQLFreeStmt(stmt, SQL_DROP);
        }
        SQLDisconnect(conn);
        SQLFreeConnect(conn);
        SQLFreeEnv(env);
    }
}
```



# Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, C++, Java, Fortran, and PL/1
- A language to which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise **embedded SQL**
- The basic form of these languages follows that of the System R embedding of SQL into PL/1
- **EXEC SQL** statement is used in the host language to identify embedded SQL request to the preprocessor  
EXEC SQL <embedded SQL statement >;
- **Note:** this varies by language
  - In some languages, like COBOL, the semicolon is replaced with END EXEC
  - In Java embedding uses #SQL { .... };

# SQLJ

- JDBC is overly dynamic, errors cannot be caught by compiler
- SQLJ: embedded SQL in Java

```
#sql iterator deptInfoIter ( String dept_name, int avgSal);
deptInfoIter iter = null;
#sql iter = { select dept_name, avg(salary)
              from instructor
              group by dept_name };
while (iter.next()) {
    String deptName = iter.dept_name();
    int avgSal = iter.avgSal();
    System.out.println(deptName + " " + avgSal);
}
iter.close();
```



## Embedded SQL

- Before executing any SQL statements, the program must first connect to the database
  - EXEC SQL **connect to** *server* **user** *user-name* **using** *password*;
- Variables of the host language can be used within embedded SQL statements
  - They are preceded by a colon (:) to them distinguish from SQL variables
    - E.g., *:credit\_amount*
- Variables must be declared within the DECLARE section
  - The syntax for declaring the variables follows the host language syntax
    - EXEC SQL BEGIN DECLARE SECTION;
    - int credit\_amount* ;
    - EXEC SQL END DECLARE SECTION;

# Embedded SQL

```
int main() {
    EXEC SQL INCLUDE SQLCA;
    EXEC SQL BEGIN DECLARE SECTION;
        int OrderID;          /* Employee ID (from user) */
        int CustID;           /* Retrieved customer ID */
        char SalesPerson[10] /* Retrieved salesperson name */
        char Status[6]        /* Retrieved order status */
    EXEC SQL END DECLARE SECTION;

    /* Set up error processing */
    EXEC SQL WHENEVER SQLERROR GOTO query_error;
    EXEC SQL WHENEVER NOT FOUND GOTO bad_number;

    /* Prompt the user for order number */
    printf ("Enter order number: ");
    scanf_s("%d", &OrderID);

    /* Execute the SQL query */
    EXEC SQL SELECT CustID, SalesPerson, Status
        FROM Orders
        WHERE OrderID = :OrderID
        INTO :CustID, :SalesPerson, :Status;

    /* Display the results */
    printf ("Customer number: %d\n", CustID);
    printf ("Salesperson: %s\n", SalesPerson);
    printf ("Status: %s\n", Status);
    exit();

query_error:
    printf ("SQL error: %ld\n", sqlca->sqlcode);
    exit();

bad_number:
    printf ("Invalid order number\n");
    exit();
}
```



## Embedded SQL

- To write an embedded SQL query, use  
**declare c cursor for <SQL query>**    *-- c used to identify the query*
- E.g.,
  - From within a host language, find the ID and name of students who have completed more than the number of credits stored in variable **credit\_amount** in the host language
  - Specify the query in SQL as follows:

EXEC SQL

```
declare c cursor for  
select ID, name  
from student  
where tot_cred > :credit_amount
```

END-EXEC

## Embedded SQL

- A variable called SQLSTATE in the SQL communication area (SQLCA) gets set to '02000' to indicate no more data is available
- The **close** statement causes the database system to delete the temporary relation that holds the result of the query

EXEC SQL **close** c ;

- **Note:** above details vary with language
  - E.g., Java embedding defines Java iterators to step through result tuples



## Updated through Embedded SQL

- Embedded SQL expressions for database modification
  - **update**, **insert**, and **delete**
- Can update tuples fetched by cursor by declaring that the cursor is for update  
EXEC SQL

```
declare c cursor for  
select *  
from instructor  
where dept_name = 'Music'  
for update
```

- Iterate through the tuples by performing **fetch** operations on the cursor, and after fetching each tuple we execute the following code:

```
update instructor  
set salary = salary + 1000  
where current of c
```

# Acknowledgements

- tutorialspoint
  - <https://www.tutorialspoint.com/what-is-the-use-of-the-method-setautocommit-in-jdbc>
- Psycopg – PostgreSQL database adapter for Python
  - <https://www.psycopg.org/docs/cursor.html>